
Beyond Distance: Geometric Competition Explains Ambiguity in Neural and LLM Representations

Anonymous Author(s)

Affiliation

Address

email

Abstract

A widely used geometric heuristic is that uncertainty is governed by radial distance to class prototypes: points near centroids are treated as certain, while distant points are treated as ambiguous. We show this account is incomplete. Studying emotion representations across eight large language model (LLM) variants and fMRI from 40 human subjects, we find that both artificial and biological systems organise emotional content in a consistent radial structure — a dense shell of observed instances at intermediate distances from class centroids, with reduced density in the immediate vicinity of the centroids. Critically, ambiguity within this structure is not determined by raw distance to a prototype but by relative competition between prototypes. We formalise this as centroid margin and demonstrate that it predicts classifier uncertainty in brain representations substantially better than raw distance, with the relationship collapsing under feature shuffling. In fine-tuned LLMs, geometric margin is tightly coupled to logit confidence. We further find that emotional content is encoded in a low-rank linear subspace occupying a small fraction of embedding dimensions, with additional dimensions contributing limited signal. Despite this shared local geometry, global representational organisation is inverted: LLMs are structured primarily along valence while brain representations are structured primarily along arousal, producing near-opposite dissimilarity matrices that strengthen as biological noise is reduced. These findings establish a reusable geometric framework for comparing representational ambiguity across artificial and neural systems, while revealing a fundamental divergence in how the two systems prioritise the affective dimensions of emotional space.

1 Introduction

How emotion is represented geometrically in artificial and biological neural systems is a foundational question at the intersection of representation learning and affective neuroscience. While large language models have achieved strong performance on emotion classification, the internal geometry supporting these classifications remains poorly characterised. The fMRI literature has documented that emotions are encoded in distributed cortical patterns [Horikawa et al., 2020], but the geometric relationship between those patterns and the structured embedding spaces of language models has not been systematically investigated.

This paper addresses a specific and tractable version of this question: do LLM embedding spaces and human fMRI activation spaces share a geometric principle for representational ambiguity? We operationalise ambiguity through *centroid distance* and *centroid margin* — the competitive gap between a point’s distance to its own class prototype versus the nearest competing prototype. Our hypothesis is that both systems exhibit a structured void near class prototypes, a peak-density belt

of real instances at moderate distances, and a gradient from certainty to ambiguity as geometric exposure to competing prototypes increases.

Research gap. Prior work on LLM–brain alignment has focused on language-processing tasks using RSA or encoding models [Caucheteux and King, 2022, Schrimpf et al., 2021, Li et al., 2023b, Huth et al., 2016, Li et al., 2023a], relating brain regions to model layers. Emotion representation geometry has been studied in affective computing [Russell, 1980, Koelstra et al., 2012] but rarely with the aim of characterising the full distributional structure of embedding spaces across a void–belt–margin framework. Whether the ambiguity gradient is a shared principle across artificial and biological systems has not been previously investigated.

Contributions.

1. A normalised-distance framework enabling geometric comparison across spaces of different dimensionality, identifying the Void and the Belt as consistent structural features across all tested systems.
2. Characterisation of overlap rates and their pairwise structure across eight LLM variants.
3. **The Low-Rank Emotion Subspace:** emotion is a low-rank linear property occupying $\approx 2.6\%$ of embedding capacity (20 of 768 dimensions). Accuracy in 20D equals or exceeds 768D across all tested architectures.
4. A margin-vs-distance diagnostic distinguishing relational from radial geometric structure: centroid *margin* predicts brain uncertainty (AUC = 0.81); centroid *distance* fails (AUC = 0.56).
5. The global representational inversion: LLMs sort emotions by valence; the brain by arousal. This produces near-opposite dissimilarity matrices that deepen as biological noise is reduced.

2 Related Work

Representational geometry in LLMs. Mikolov et al. [2013] showed that distributional embeddings encode semantic structure in linear subspaces. Ethayarajh [2019] documented anisotropy in contextualised representations. Kumar et al. [2022] showed fine-tuning can compress pretrained features — directly relevant to our compaction finding.

Brain–LLM alignment. Caucheteux and King [2022] showed partial convergence between brain activations and LLMs during language processing. Schrimpf et al. [2021] characterised neural architecture convergence with transformer predictive processing. Huth et al. [2016] linked distributional semantics to cortical topography. Ismayilzada et al. [2026] found alignment during creative thinking (preprint).

Distributed neural coding of emotion. Horikawa et al. [2020] showed visually evoked emotion representations are high-dimensional, categorical, and distributed across transmodal brain regions. Haxby et al. [2001] established that cognitive categories are encoded as distributed activation patterns — the biological prior explaining why brain fMRI shows negative silhouette in 2D projection while remaining decodable in 48 dimensions.

The circumplex model. Russell [1980] organised affective states along valence and arousal dimensions, operationalised through DEAP [Koelstra et al., 2012], ANEW [Bradley and Lang, 1999], and IAPS [Lang et al., 2008]. We use external V-A reference coordinates from this literature to interpret global representational axes, not as measured participant ground truth.

3 Datasets and Models

Text emotion dataset. We use the `dair-ai/emotion` dataset [Saravia et al., 2018], a 6-class text emotion corpus (anger, fear, happiness, love, sadness, surprise). We use the validation split, balanced to **4,038 samples** (673 per class). Class balance ensures centroids are not displaced by class-size differences and that overlap rates are comparable across categories. Labels are integer-encoded: anger=0, fear=1, happiness=2, love=3, sadness=4, surprise=5.

84 **Brain fMRI dataset.** We use OpenNeuro DS005700 [OpenNeuro DS005700, 2024], “Neural MO
85 — fMRI Dataset for Emotion Recognition.” The dataset contains **40 subjects** viewing emotion-
86 eliciting stimuli under five conditions: afraid, calm, delighted, depressed, and excited. Each subject
87 contributes activations across **48 cortical ROIs** from the Harvard-Oxford atlas, yielding 200 obser-
88 vations total (40 subjects $\times$ 5 emotions). Repeated blocks per subject and emotion are averaged
89 to stable ROI vectors. Brain LOSO decoding accuracy is 0.56 (95% CI: 0.49–0.63, permutation
90 $p < 0.001$), well above the 5-class chance level of 0.20. Our $N = 40$ exceeds high-impact prece-
91 dents with smaller samples: Individual Brain Charting ($N = 12$) [OpenNeuro DS005700, 2024],
92 Precision Functional Mapping of Children ($N = 12$), and Natural Scenes Dataset ($N = 8$).

93 **LLM variants.** We evaluate eight variants by crossing two base models, two training states, and
94 two extraction layers:

Variant	Base Model	State	Layer	Dim
MPNet-FT-Final	all-mpnet-base-v2 [Song et al., 2020]	Fine-tuned	L12	768
MPNet-FT-Mid	all-mpnet-base-v2	Fine-tuned	L6	768
MPNet-Base-Final	all-mpnet-base-v2	Pretrained	L12	768
MPNet-Base-Mid	all-mpnet-base-v2	Pretrained	L6	768
BGE-FT-Final	bge-base-en-v1.5 [Xiao et al., 2023]	Fine-tuned	L12	768
BGE-FT-Mid	bge-base-en-v1.5	Fine-tuned	L6	768
BGE-Base-Final	bge-base-en-v1.5	Pretrained	L12	768
BGE-Base-Mid	bge-base-en-v1.5	Pretrained	L6	768

96 Fine-tuned models were trained on the `dair-ai/emotion` training split. The cross-system context
97 retention experiment (Section 5.2) additionally includes Qwen3-1.7B fine-tuned [Qwen Team, Al-
98 ibaba Cloud, 2025] (PCA-projected to 768D) as a third architecture to test whether the compaction-
99 distribution finding generalises beyond the two primary model families.

100 4 Methods

101 **Embedding extraction.** All text embeddings use attention-mask-weighted mean pooling:

$$\mathbf{e} = \frac{\sum_{i=1}^N \mathbf{h}_i \cdot m_i}{\sum_{i=1}^N m_i} \quad (1)$$

102 where $\mathbf{h}_i$ is the hidden state for token i and m_i is the attention mask. Final-layer
103 (L12) extraction uses the last hidden state directly. Middle-layer (L6) extraction enables
104 `output_hidden_states=True` and accesses `hidden_states[6]`. Tokenisation: model-specific
105 tokenisers with `max_length=128`, `padding=True`, `truncation=True`, batch size 32.

106 **Normalised distance framework.** To compare across systems with different intrinsic scales (768D
107 LLMs vs. 48D brain ROI), all distances are normalised by the mean radius of each system:

$$\tilde{d}(x, c_k) = \frac{d(x, c_k)}{\bar{R}_k}, \quad \bar{R}_k = \frac{1}{N_k} \sum_{x \in C_k} d(x, c_k) \quad (2)$$

108 This maps all systems to a common scale (≈ 0 –2.5 normalised units) without discarding distri-
109 butional shape. A point geometrically overlaps class C' (its non-assigned class) if $d(P, c_{C'}) <$
110 $d(P, c_C)$. The *centroid margin* is:

$$m(P) = d(P, c_{\text{nearest competitor}}) - d(P, c_C) \quad (3)$$

111 Positive margin: P is closer to its true centroid than to any competitor (more certain). Negative
112 margin: geometric overlap (more ambiguous).

113 **Brain 48D→11D transformation.** To test whether the global representational inversion (Sec-
114 tion 5.5) is noise-driven, we construct a domain-informed 11-dimensional representation by aggre-
115 gating ROI activations into biologically coherent groups. The Default Mode Network (DMN) [Li
116 et al., 2014] supports self-referential processing and social cognition. The Salience Network [Seeley,
117 2019] mediates detection of affectively significant stimuli. The Central Executive Network [Miller

and Cohen, 2001] coordinates regulatory responses to emotional content. The 11 dimensions consist of five anatomical lobe means (frontal, temporal, parietal, occipital, cingulate/limbic), five functional network means (DMN, Salience, Central Executive, visual, somatomotor), and one neighbourhood feature capturing local inter-regional coordination. This is domain-informed structured denoising, not generic PCA: each dimension corresponds to a biologically meaningful functional unit. Classification performance in 11D remains above chance and comparable to 48D, confirming the transformation preserves task-relevant signal.

Valence-arousal reference mapping. External V-A coordinates are drawn from the circumplex model [Russell, 1980, Mehrabian and Russell, 1974] and affective-norm resources DEAP [Koelstra et al., 2012], ANEW [Bradley and Lang, 1999], and IAPS [Lang et al., 2008] (1–9 SAM scale). Because the fMRI dataset supplies categorical labels rather than participant-level V-A ratings, these coordinates function as an *external reference geometry*, not measured subjective ground truth. Sensitivity analysis (Appendix A) perturbs coordinates with Gaussian noise (SD=0.25/0.50/0.75) to verify qualitative robustness.

RSA methodology. We follow Kriegeskorte [2008]. For each system, pairwise Euclidean distances between emotion class centroids form a $K \times K$ Representational Dissimilarity Matrix (RDM). Brain-LLM comparison uses the 3-emotion overlap: fear $\approx$ afraid, happiness $\approx$ delighted, sadness $\approx$ depressed. Systems are compared by Pearson correlation between vectorised RDMs. A brain-brain noise ceiling estimates the upper bound imposed by individual subject variability: upper 0.484, lower 0.460 (mean ≈ 0.47), computed by split-half resampling of the 40-subject dataset. With $K = 3$ emotions, RDMs have 1 degree of freedom ($p \approx 0.21$ for $r = -0.88$; formally non-significant); load-bearing quantities are sign and noise-ceiling contrast. The 11D result (Section 5.5) provides stronger support via a fully negative bootstrap CI.

5 Results


We present five findings in order of increasing cross-system scope. Findings 1–3 characterise LLM geometry; Finding 4 introduces the ambiguity gradient; Finding 5 documents the global representational inversion.

5.1 Finding 1: The Void and the Belt

A consistent two-zone density structure emerges across all 8 LLM variants and human brain fMRI (Figure 1). The structural parallel is qualitative: packing density magnitudes differ substantially across systems, and density-shape similarity does not reach conventional significance under permutation testing (Wasserstein $p = 0.46$). The claim is descriptive — shared organisational logic, not matched compression.

The Void. No data points exist within approximately 0.375–0.8125 normalised distance units of any emotion centroid. Fine-tuned models show first non-zero density at ≈ 0.4375 units; pretrained models at ≈ 0.6875 (BGE-Base) and ≈ 0.8125 (MPNet-Base) units. The Void arises as a geometric consequence of distributed within-class variance: a centroid is the mean of a distributed ring of instances, so it sits inside the ring in zero-density space.

The Belt. Density rises sharply from the void boundary and peaks at approximately 0.56–0.94 normalised units. Fine-tuned models peak earlier (MPNet-FT-Final at 0.5625, BGE-FT-Final at 0.6875); pretrained models peak later (both at 0.9375).



(a) Radial density profiles (top) and ambiguity gradients (bottom) for all 8 LLM variants. The void region (zero density near centroid) and the Belt (peak density shell) are consistent across all systems.

(b) Radial density decay for LLM variants (coloured) overlaid with human brain fMRI (dashed black). The two-zone structure extends to neural data.

Figure 1: The Void and the Belt (Finding 1). (a) All 8 LLM variants share the same two-zone radial density structure in normalised-distance space. (b) The same structure is present in human brain fMRI emotion representations, establishing it as a cross-system geometric principle.

Certainty buffer. The radial gap between density peak and onset of geometric ambiguity ($> 5\%$ overlap rate; geometric ambiguity defined as the region where nearest-neighbour class boundaries overlap in embedding space):

Variant	Void Onset	Density Peak	Overlap Onset	Certainty Buffer
MPNet-FT-Final	0.4375	0.562	2.438	+1.875
BGE-FT-Final	0.4375	0.688	2.438	+1.750
MPNet-Base-Final	0.8125	0.938	0.938	0.000
BGE-Base-Final	0.6875	0.938	1.062	+0.125

Fine-tuning creates a geometrically safe zone between peak density and the onset of competitive class confusion; for pretrained models this buffer is near zero. No embedding aligns exclusively with a single class prototype in any system: global overlap rates range from 1.76% (MPNet-FT-Final) to 19.22% (BGE-Base-Final), confirming that fine-tuning substantially reduces inter-class geometric confusion.


5.2 Finding 2: Information Compaction vs. Distribution

Iterative SVD ablation removes the most informative linear direction at each step; a fresh logistic regression is re-trained on the residual. **Signal erasure points** (dimension at which accuracy reaches the 6-class chance baseline of 16.7%):

Variant	Baseline Accuracy (%)	Erasure Point
MPNet-FT-Final	98.24	dim 15
BGE-FT-Final	97.99	dim 20
BGE-Base-Final	93.78	dim 26
MPNet-Base-Final	94.08	dim 67

Fine-tuned models achieve higher baseline accuracy but concentrate all emotional signal into 15–20 directions; remove them and signal collapses entirely. Pretrained MPNet-Base retains decodable signal for 67 dimensions, a $4.5\times$ wider manifold than the fine-tuned equivalent (Figure 2a).

Cross-system extension. A separate iterative ablation across three systems (Figure 2b) shows qualitatively distinct signal geometries: the human brain ($D_{50} = 4$, AUC 0.302) uses the most distributed code — removing just 4 dominant directions halves accuracy, yet the curve flattens into a plateau rather than collapsing to chance [Haxby et al., 2001]; Qwen-FT ($D_{50} = 12$, AUC 0.260) compresses



(a) Signal cliff: accuracy as dominant directions are removed iteratively. Fine-tuned models reach chance at dimensions 15–20; pretrained BGE at 26, pretrained MPNet at 67. (b) Cross-system context retention. Brain $D_{50} = 4$; fine-tuned LLM $D_{50} = 12$; pretrained LLM $D_{50} = 17$.

Figure 2: **Signal structure across systems (Finding 2).** (a) SVD ablation reveals qualitatively different compaction strategies. (b) The brain requires only 4 dominant directions to achieve half its classification accuracy, yet maintains a distributed plateau rather than collapsing to chance — consistent with distributed biological coding.

184 signal into a steep cliff; pretrained MPNet ($D_{50} = 17$, AUC 0.331) spreads it across the widest
185 manifold.

186 5.3 Finding 3: The Low-Rank Emotion Subspace

187 **Named result.** Emotion is a low-rank linear property, occupying $\approx 20/768 \approx 2.6\%$ of full embed-
188 ding capacity. The cloud appearance of pretrained models in 768D is not an absence of geometry —
189 it is the geometry of ≈ 748 noise dimensions overwhelming 20 signal dimensions.

190 **Silhouette scores and classification accuracy before and after projection into the top 20 SVD**
191 **directions (Figure 3):**

Variant	768D Sil.	20D Sil.	Sil. Change	768D Acc. (%)	20D Acc. (%)
MPNet-FT-Final	0.860	0.907	+5.5%	98.24	98.27
BGE-FT-Final	0.685	0.835	+22.0%	97.99	98.09
MPNet-Base-Final	0.047	0.408	+765%	94.08	95.64
BGE-Base-Final	0.057	0.415	+627%	93.78	95.12

193 The extra 748 dimensions contribute only noise that slightly suppresses linear classification per-
194 formance. Pairwise centroid distances preserve the same relational structure in 20D as in 768D:
195 Happiness–Love and Fear–Anger are the closest pairs; Sadness–Surprise and Happiness–Fear are
196 the most distant.

197 5.4 Finding 4: The Cross-System Ambiguity Gradient

198 Both systems show that geometric positioning relative to class prototypes predicts classification
199 confidence — through mechanistically distinct constructs.

200 **LLMs: centroid margin $\rightarrow$ logit margin.** For fine-tuned LLMs, centroid margin is a strong pre-
201 dictor of logit margin (Figure 4a):

Variant	Margin-logit Pearson r	Logit agreement in overlap (%)
BGE-FT-Final	0.9572	93.4
MPNet-FT-Final	0.9884	100.0

203 Fine-tuning aligns geometric and logit structure by design, so a positive correlation is expected.
204 What is not guaranteed by the training objective is the *degree*: a near-linear relationship ($r = 0.9884$)
205 means centroid margin almost perfectly determines logit margin, with no residual variance at-
206 tributable to non-geometric factors. In overlap regions specifically, logit values are biased toward the





Figure 3: **The Low-Rank Emotion Subspace (Finding 3).** Silhouette scores in full 768D (grey) vs. the isolated top-20 SVD subspace (red). Pretrained models improve from near-zero to ≈ 0.41 ; fine-tuned models improve modestly from already-high baselines. Accuracy in 20D equals or exceeds 768D in all four variants (see text). PCA scatter visualisations are in Appendix F.

207 competing class 93.4–100% of the time — meaning geometry at the decision boundary is fully de-
 208 terminative of classifier confidence, not merely correlated with it. This result is reported as a mecha-
 209 nistic characterisation, not a novel discovery: its value is establishing that the geometric competition
 210 framework (centroid margin) is the operative construct in LLMs, which motivates its application to
 211 the brain in the following section.

212 **Brain: centroid margin $\rightarrow$ classifier uncertainty.** Classifier uncertainty is operationalised as $1 -$
 213 $P(\text{correct})$, the complement of the LOSO classifier’s predicted probability for the true emotion
 214 class. For AUC computation, samples are binarised at the median: *high uncertainty* is defined
 215 as $1 - P(\text{correct}) > \text{median}$. For brain fMRI, raw centroid distance is a weak predictor of this
 216 uncertainty (Spearman $r = 0.1509$, $\text{AUC} = 0.5627$ — near chance). The correct predictor is centroid
 217 margin (Equation 3):

$$\text{AUC}_{\text{margin}} = 0.8124 \quad \text{vs.} \quad \text{AUC}_{\text{raw distance}} = 0.5627$$

218 with Spearman $r = 0.6108$, $p = 7.78 \times 10^{-22}$. A feature-shuffle control ($r = 0.0287$, $p = 0.687$)
 219 confirms the relationship vanishes under random permutation.



(a) LLM: centroid margin vs. logit margin for MPNet-FT-Final ($r = 0.988$). Geometric margin almost perfectly determines logit margin.

(b) Brain fMRI: centroid margin predicts classifier uncertainty ($\text{AUC} = 0.8124$); raw centroid distance fails ($\text{AUC} = 0.5627$). Margin is the operative geometric construct.

Figure 4: **Geometric predictors of confidence (Finding 4).** (a) In LLMs, centroid margin tightly predicts logit margin. (b) In the brain, centroid *margin* — not distance — predicts uncertainty, revealing competition between prototypes as the key signal.

220 **The ambiguity gradient comparison.** The ambiguity-gradient curves exhibit a consistent mono-
 221 tonic structure across systems, with ambiguity increasing as a function of normalised distance from
 222 class centroids. This qualitative alignment is supported by a full-pipeline permutation test ($n = 5000$
 223 shuffles, $p < 0.001$) that destroys cross-system correspondence under label shuffling. Correlation

coefficients computed over interpolated grid points are not treated as inferential statistics, as adjacent points are not independent observations; the permutation test is the sole inferential basis for this comparison. KS statistic = 0.57 and Cohen’s $d = 1.34$ confirm that while gradient *shape* is shared, absolute packing density differs substantially between systems — shared ambiguity structure, distinct representational compaction.

5.5 Finding 5: The Valence-Arousal Inversion

The most novel finding of the study is that the shared local ambiguity principle coexists with a global inversion of representational geometry.

(5a) RDM comparison in 48D raw ROI space (Appendix H).

Comparison	r (direction only) [†]	Notes
Brain vs. fine-tuned LLM (cosine)	−0.8756	Directional evidence of global inversion
Brain vs. pretrained LLM (cosine)	−0.1023	Confirmed negative
Brain vs. fine-tuned LLM (Manhattan)	−0.5476	Metric-invariant validation
Brain-brain noise ceiling	+0.47	Upper 0.484, lower 0.460

[†] $K = 3$, 1 df; $p \approx 0.21$ for $r = -0.88$; no inferential p-value is valid.

RDM comparisons provide directional evidence of global inversion, but because only three shared emotion categories are available, exact correlation magnitudes are unstable (bootstrap 95% CI: $[-0.9967, +0.6994]$, as expected with $K = 3$, 1 df; the CI spans positive values and cannot support directional inference at this dimensionality). Brain subjects agree at $r \approx 0.47$, yet the fine-tuned LLM RDM points in the *opposite direction*; sign and noise-ceiling contrast are load-bearing. The 11D bootstrap CI $[-0.9899, -0.1178]$, fully negative) provides stronger directional support (Section 5.5).

(5b) RDM comparison in 11D systems-level space.

Comparison	48D r	11D r
Brain vs. fine-tuned LLM	−0.6371	−0.7539
Brain vs. pretrained LLM	−0.1023	−0.9918

As biological signal is denoised from 48 ROIs to 11 systems-level features, the divergence *amplifies* rather than shrinks. The pretrained LLM at 11D ($r = -0.9918$) is the strongest cross-system result in the study; this specific magnitude carries the uncertainty of a 3-point Pearson correlation (1 df) and should be treated as indicative. The robust directional claim is: as biological noise is removed, the negative correlation between brain and LLM representational geometry becomes more extreme, not less — inconsistent with the hypothesis that the paradox is noise-driven.


(5c) Valence-arousal axis alignment. We projected emotion centroids into the leading 2D PCA subspace and correlated each axis with external V-A reference coordinates (Table 1).

251

Emotion	Valence	Arousal	Quadrant
Afraid / Fear	−0.80	0.85	negative, high arousal
Calm	+0.70	0.15	positive, low arousal
Delighted / Happin.	+0.85	0.70	positive, moderate–high arousal
Depressed / Sadness	−0.85	0.25	negative, low arousal
Excited	+0.75	0.90	positive, high arousal
Love	+0.85	0.55	positive, moderate arousal

Table 1: External V-A reference coordinates (normalised circumplex scale: valence −1 to +1, arousal 0 to 1). The quadrant structure — not exact decimal values — is the scientifically load-bearing element.

255



(a) Brain fMRI: poor VA alignment (Disparity=0.80, $R_{\text{val}} = 0.31$). Embeddings (red crosses) deviate markedly from the V-A circumplex (blue circles). (b) MPNet-Balanced (fine-tuned): excellent valence alignment (Disparity= 0.19, $R_{\text{val}} = 0.97$). Embeddings closely track the horizontal valence axis.

Figure 5: **The Valence-Arousal Inversion (Finding 5c).** LLMs organise emotion primarily by valence; the brain primarily by arousal. The global representational priorities are swapped.

System	PC1 axis	PC1 r	PC2 axis	Permutation p
Fine-tuned LLM	Valence	0.98	Arousal ($r=0.82$)	0.013
Pretrained LLM	Valence	0.97	Arousal ($r=0.73$)	0.009
Human Brain	Arousal	0.96	Valence ($r=0.31$)	0.033

Table 2: Global representational axis alignment (Figure 5). The axis priorities are swapped between LLMs and the brain.

Both LLM types are valence-dominant on PC1 ($r = 0.97$ – 0.98); the brain is arousal-dominant ($r = 0.96$), with only weak valence alignment on PC2 ($r = 0.31$). Because K is small and V-A coordinates are category-level approximations, these results are global-structure evidence, not primary proof: the dominant neural axis aligns more strongly with arousal than valence, while the ambiguity mechanism is supported independently by centroid margin and classifier uncertainty.

6 Discussion

The shared geometric principle. Both LLM embedding spaces and the tested brain fMRI dataset exhibit the same general relationship: geometric positioning relative to class prototypes predicts representational certainty. In LLMs this operates via centroid margin predicting logit margin; in the brain, via centroid margin predicting classifier uncertainty. These are complementary implementations of a geometric competition principle. The mechanistic interpretation likely lies in the mathematics of distributed high-dimensional representations: any system encoding discrete categories via distributed activation will produce centroids surrounded by distributed instances, generating a void, a belt, and an ambiguity gradient.

Statistical hierarchy of evidence. The findings differ in evidentiary strength. The *strongest* evidence is the margin-uncertainty coupling: brain centroid margin predicts classifier uncertainty (AUC = 0.81 vs. 0.56 for raw distance), collapsing under feature shuffling; LLM centroid margin predicts logit margin ($r = 0.99$, permutation $p = 0.00$). *Strong* evidence: the valence–arousal axis inversion, supported by independent permutation tests on PC-axis correlations ($p < 0.05$ for all three systems). *Supportive* evidence: the 11D RDM inversion, whose bootstrap CI for fine-tuned LLMs is fully negative. *Descriptive evidence only:* the 48D RDM magnitude (3-point Pearson, wide CI) and the void–belt density similarity (Wasserstein permutation $p = 0.46$). The primary evidence for global inversion is therefore the valence–arousal axis analysis, not the 3-emotion RSA magnitude.

The relational paradox: different priorities, not failure. The inverted RDM directionality reflects different operational demands, not error. LLMs encode co-occurrence statistics: “fear” and “sadness” share negative hedonic contexts, clustering centroids along valence. The brain evolved to

respond to threats and rewards: fear’s high-arousal physiology differs fundamentally from sadness’s low-arousal withdrawal, so arousal is the more behaviourally predictive axis.

Limitations. Our evidence covers one text emotion dataset (6 classes) and one brain fMRI dataset (5 emotions, 40 subjects, one ROI atlas). The 3-emotion cross-system label mapping (fear≈afraid, happiness≈delighted, sadness≈depressed) relies on conceptual equivalence and introduces non-trivial uncertainty into all RSA comparisons. The V-A analysis relies on externally assigned category-level coordinates rather than participant-specific ratings; the global-axis findings should be interpreted as supportive evidence, not as the main proof of the margin-based ambiguity mechanism. All findings are observational; no causal evidence is established.

7 Conclusion

We have presented five findings characterising the geometry of emotion representations in LLMs and human fMRI. The core result is a paradox: a geometric competition principle for representational ambiguity is observable across both systems — LLMs via centroid margin predicting logit margin; the brain via centroid margin predicting classifier uncertainty — yet global representational organisation is inverted: LLMs sort emotions by valence, the brain by arousal.

Three findings have practical value: the low-rank subspace (2.6% capacity, 20D) shows efficient emotion encoding does not require large representations; the margin-vs-distance contrast (AUC 0.81 vs. 0.56) provides a reusable diagnostic; and the V-A inversion suggests training on arousal distinctions may be necessary for biological alignment.

More broadly, the coexistence of shared local geometry with globally inverted structure suggests that cross-system alignment in affective representation is not a binary question: two systems can share the same local competitive logic while differing fundamentally in what they prioritise globally.

References

- E. J. Allen, G. St-Yves, Y. Wu, J. L. Breedlove, J. S. Prince, L. T. Dowdle, M. Nau, B. Caron, F. Pestilli, I. Charest, J. B. Hutchinson, T. Naselaris, and K. Kay. A massive 7T fMRI dataset to bridge cognitive neuroscience and artificial intelligence. *Nature Neuroscience*, 25(1):116–126, 2022. doi: 10.1038/s41593-021-00962-x.
- M. M. Bradley and P. J. Lang. Affective norms for English words (ANEW): Instruction manual and affective ratings. Technical Report Technical Report C-1, University of Florida, Center for Research in Psychophysiology, 1999.
- C. Caucheteux and J.-R. King. Brains and algorithms partially converge in natural language processing. *Communications Biology*, 5(1), 2022. doi: 10.1038/s42003-022-03036-1.
- K. Ethayarajh. How contextual are contextualized word representations? comparing the geometry of BERT, ELMo, and GPT-2 embeddings. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019. doi: 10.18653/v1/D19-1006.
- E. M. Gordon, T. O. Laumann, A. W. Gilmore, D. J. Newbold, D. J. Greene, J. J. Berg, M. Ortega, C. Hoyt-Drazen, C. Gratton, H. Sun, J. M. Hampton, R. S. Coalson, A. L. Nguyen, K. B. McDermott, J. S. Shimony, A. Z. Snyder, B. L. Schlaggar, S. E. Petersen, S. M. Nelson, and N. U. F. Dosenbach. Precision functional mapping of individual human brains. *Neuron*, 95(4):791–807, 2017. doi: 10.1016/j.neuron.2017.07.011.
- J. V. Haxby, M. I. Gobbini, M. L. Furey, A. Ishai, J. L. Schouten, and P. Pietrini. Distributed and overlapping representations of faces and objects in ventral temporal cortex. *Science*, 293(5539):2425–2430, 2001. doi: 10.1126/science.1063736.
- T. Horikawa, A. S. Cowen, D. Keltner, and Y. Kamitani. The neural representation of visually evoked emotion is high-dimensional, categorical, and distributed across transmodal brain regions. *iScience*, page 101060, 2020. doi: 10.1016/j.isci.2020.101060.
- A. G. Huth, W. A. de Heer, T. L. Griffiths, F. E. Theunissen, and J. L. Gallant. Natural speech reveals the semantic maps that tile human cerebral cortex. *Nature*, 532(7600):453–458, 2016. doi: 10.1038/nature17637.

334 M. Ismayilzada et al. Large language models align with the human brain during creative thinking.
335 *arXiv preprint arXiv:2604.03480*, 2026.

336 S. Koelstra, C. Mühl, M. Soleymani, J.-S. Lee, A. Yazdani, T. Ebrahimi, T. Pun, A. Nijholt, and
337 I. Patras. DEAP: A database for emotion analysis using physiological signals. *IEEE Transactions*
338 *on Affective Computing*, 3(1):18–31, 2012. doi: 10.1109/T-AFFC.2011.15.

339 N. Kriegeskorte. Representational similarity analysis — connecting the branches of systems neuro-
340 science. *Frontiers in Systems Neuroscience*, 2, 2008. doi: 10.3389/neuro.06.004.2008.

341 A. Kumar, A. Raghunathan, R. Jones, T. Ma, and P. Liang. Fine-tuning can distort pretrained features
342 and undermine catastrophic forgetting. In *International Conference on Learning Representations*
343 *(ICLR)*, 2022.

344 P. J. Lang, M. M. Bradley, and B. N. Cuthbert. International affective picture system (IAPS): Affec-
345 tive ratings of pictures and instruction manual. Technical Report Technical Report A-8, University
346 of Florida, NIMH Center for the Study of Emotion and Attention, 2008.

347 J. Li, A. Karamolegkou, Y. Kementchedjhieva, M. Abdou, S. Lehmann, and A. Søgaard. Structural
348 similarities between language models and neural response measurements. In *Proceedings of the*
349 *37th Annual Conference on Neural Information Processing Systems (NeurIPS)*, pages 346–365,
350 2023a.

351 J. Li, A. Karamolegkou, Y. Kementchedjhieva, and A. Søgaard. Large language models converge
352 on brain-like word representations. *arXiv preprint arXiv:2306.01930*, 2023b.

353 W. Li, X. Mai, and C. Liu. The default mode network and social understanding of others: what
354 do brain connectivity studies tell us. *Frontiers in Human Neuroscience*, 8, 2014. doi: 10.3389/
355 fnhum.2014.00074.

356 A. Mehrabian and J. A. Russell. *An Approach to Environmental Psychology*. MIT Press, Cambridge,
357 MA, 1974.

358 T. Mikolov, K. Chen, G. Corrado, and J. Dean. Distributed representations of words and phrases and
359 their compositionality. In *Advances in Neural Information Processing Systems*, volume 26, 2013.

360 E. K. Miller and J. D. Cohen. An integrative theory of prefrontal cortex function. *Annual Review of*
361 *Neuroscience*, 24:167–202, 2001. doi: 10.1146/annurev.neuro.24.1.167.

362 OpenNeuro DS005700. Neural MO — fMRI dataset for emotion recognition. OpenNeuro. <https://openneuro.org/datasets/ds005700>, 2024.

364 A. L. Pinho, A. Amadon, T. Ruest, M. Fabre, E. Dohmatob, I. Denghien, C. Ginisty, S. Becuwe-
365 Desmidt, S. Roger, and B. Thirion. Individual brain charting, a high-resolution fMRI dataset for
366 cognitive mapping. *Scientific Data*, 5:180105, 2018. doi: 10.1038/sdata.2018.105.

367 J. Posner, J. A. Russell, and B. S. Peterson. The circumplex model of affect: An integrative approach
368 to affective neuroscience, cognitive development, and psychopathology. *Development and Psy-*
369 *chopathology*, 17(3):715–734, 2005. doi: 10.1017/S0954579405050340.

370 Qwen Team, Alibaba Cloud. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

371 J. A. Russell. A circumplex model of affect. *Journal of Personality and Social Psychology*, 39(6):
372 1161–1178, 1980. doi: 10.1037/h0077714.

373 E. Saravia, H.-C. T. Liu, Y.-H. Huang, J. Wu, and Y.-S. Chen. CARER: Contextualized affect
374 representations for emotion recognition. In *Proceedings of the 2018 Conference on Empirical*
375 *Methods in Natural Language Processing (EMNLP)*, pages 3687–3697, 2018. doi: 10.18653/v1/
376 D18-1404.

377 M. Schrimpf, I. A. Blank, G. Tuckute, C. Kauf, E. A. Hosseini, N. Kanwisher, J. B. Tenenbaum, and
378 E. Fedorenko. The neural architecture of language: Integrative modeling converges on predictive
379 processing. *Proceedings of the National Academy of Sciences*, 118(45):e2105646118, 2021. doi:
380 10.1073/pnas.2105646118.

- 381 W. W. Seeley. The salience network: A neural system for perceiving and responding to homeo-
382 static demands. *Journal of Neuroscience*, 39(50):9878–9882, 2019. doi: 10.1523/JNEUROSCI.
383 1138-17.2019.
- 384 K. Song, X. Tan, T. Qin, J. Lu, and T.-Y. Liu. MPNet: Masked and permuted pre-training for
385 language understanding. In *Advances in Neural Information Processing Systems*, volume 33,
386 2020.
- 387 S. Xiao, Z. Liu, P. Zhang, and N. Muennighoff. C-Pack: Packaged resources to advance general
388 Chinese embedding. *arXiv preprint arXiv:2309.07597*, 2023.

389 A V-A Reference Coordinate Robustness

390 The external V-A reference coordinates used in Section 5.5 are category-level approximations from
391 the circumplex literature [Posner et al., 2005]. To confirm that the axis-dominance results are not
392 artefacts of these specific coordinate assignments, we ran a permutation test ($n = 1000$) for each
393 system.

394 **Procedure.** For each of the three systems (fine-tuned LLM, pretrained LLM, brain fMRI), emotion
395 centroids are projected into the leading 2D PCA subspace. Valence and arousal correlations are
396 computed with the reference coordinates (Table 1). A null distribution is constructed by shuffling
397 the V-A label assignments across emotions 1,000 times; each shuffle recomputes the maximum of
398 the valence and arousal correlations. The permutation p -value is the proportion of shuffles that equal
399 or exceed the observed maximum correlation.

400 Results.

	System	Dominant axis	Observed r	Permutation p ($n = 1000$)
401	Fine-tuned LLM	Valence	0.97	0.009
	Pretrained LLM	Valence	0.97	0.013
	Brain fMRI	Arousal	0.96	0.033

402 All three systems reach conventional significance under label permutation. The axis assignments
403 — valence-dominant for both LLM types, arousal-dominant for the brain — are not recoverable by
404 random reassignment of the reference coordinates. The qualitative inversion is robust to the specific
405 coordinate values used.

406 B Phase 5 Full Results

407 Phase 5 evaluates only the two fine-tuned final-layer models against the 4,038-sample validation set.

	Metric	BGE-FT-Final	MPNet-FT-Final
	Total samples	4,038	4,038
408	Overlap samples	≈ 76 (1.88%)	≈ 71 (1.76%)
	Dist-logit Pearson r	0.9572	0.9884
	Logit agreement in overlap	93.4%	100%
	Dominant overlap pair	Happin./Love	Happin./Love

409 The 100% logit agreement for MPNet-FT-Final means every sample geometrically closer to a com-
410 peting centroid also has a higher raw logit score for that competing class. Embedding geometry is
411 the direct mechanistic basis for classifier confidence in this model.

C Overlap Metrics: All Final-Layer Variants

Variant	Overlap (%)	Density Peak	Certainty Buffer	Void Onset
MPNet-FT-Final	1.76	0.562	+1.875	≈ 0.4375
BGE-FT-Final	1.88	0.688	+1.750	≈ 0.4375
MPNet-Base-Final	18.45	0.938	0.000	≈ 0.8125
BGE-Base-Final	19.22	0.938	+0.125	≈ 0.6875
BGE-Base-Mid	28.14	—	—	—
MPNet-Base-Mid	16.92	—	—	—

D fMRI Sample Size Justification

Dataset	N subjects	Venue
Individual Brain Charting [Pinho et al., 2018]	12	Scientific Data
Precision Functional Mapping [Gordon et al., 2017]	10	Neuron
Natural Scenes Dataset [Allen et al., 2022]	8	Nature Neuroscience
This study (DS005700) [OpenNeuro DS005700, 2024]	40	OpenNeuro

With $N = 40$, LOSO cross-validation leaves 39 training subjects per fold — well within the range where logistic regression on 48 features is well-determined. The 5-class chance level is 0.20; our LOSO accuracy of 0.56 (95% CI: 0.49–0.63, permutation $p < 0.001$) represents an absolute lift of 0.36 above chance.


E 48D→11D Transformation: ROI-to-Network Mapping

Anatomical lobe features (5): Mean activation within (1) frontal cortex ROIs, (2) temporal cortex ROIs, (3) parietal cortex ROIs, (4) occipital cortex ROIs, (5) cingulate/limbic ROIs.

Functional network features (5): Mean activation within (6) Default Mode Network ROIs [Li et al., 2014] — bilateral medial PFC, posterior cingulate, angular gyrus; (7) Salience Network ROIs [Seeley, 2019] — bilateral anterior insula, anterior cingulate cortex; (8) Central Executive Network ROIs [Miller and Cohen, 2001] — bilateral dorsolateral PFC, lateral parietal cortex; (9) visual processing network ROIs — bilateral occipital and fusiform regions; (10) somatomotor network ROIs — bilateral pre/postcentral gyrus.

Neighbour context feature (1): Mean absolute difference between adjacent ROI activations within each anatomical region, capturing local inter-regional coordination.

F Finding 3: PCA Scatter Visualisations




(a) BGE-Base-Final: PCA scatter in full 768D (left) vs. top-20 SVD subspace (right). Diffuse clouds sharpen into separated clusters. (b) MPNet-Base-Final: same comparison. Silhouette improves from 0.057 to 0.415 (+627%) and 0.047 to 0.408 (+765%).

Figure 6: **768D vs. top-20 SVD subspace (Finding 3).** Projecting into the 20-dimensional signal subspace reveals the cluster structure hidden by noise dimensions in the full embedding space.

G Image Experiments: Negative Result

An extension to 120,000 affective image embeddings (CLIP-based) was explored. Image embeddings did not yield clear cluster separation: silhouette scores remained near zero and the void-belt density structure was absent. The failure mode was insufficient cluster separation across emotion categories. This negative result informs the scope of the geometric competition principle: it is observable in text-based LLM embeddings and brain fMRI representations under the systems and datasets tested, but was not apparent in CLIP image embeddings. The modality (text vs. image), CLIP’s training objective, and the emotion taxonomy used are candidate explanations.

H RDM Heatmaps (Finding 5a)



(a) Brain RDM (cosine): arousal-structured. Depressed–calm (0.72) and afraid–excited (0.82) are close; affect separates by activation intensity.

(b) LLM RDM (cosine): valence-structured. Sadness–happiness (0.36), love–happiness (0.26) are close; affect separates by hedonic polarity.

Figure 7: **RDM comparison (Finding 5a)**. The brain organises affect by arousal intensity while the LLM organises by hedonic polarity — geometries point in opposite directions (direction indicator only; $K=3$, 1 df, $p \approx 0.21$; see table for values). Noise ceiling $\approx +0.47$.

I Compute and Reproducibility

Hardware. LLM embedding extraction: HPC cluster (GPU nodes). Brain analysis: standard compute (CPU-bound statistical analyses).

Code and data. Code and data are provided as supplementary material. The repository contains: (1) embedding extraction scripts for all 8 LLM variants, (2) all six brain analysis scripts, (3) the 48D→11D transformation code, (4) the cross-system global behaviour comparison script, (5) the V-A alignment code, (6) valence_arousal_reference_table.csv, (7) requirements and environment specification.

Ethics statement. This study uses only publicly available, de-identified neuroimaging data from the Neural MO dataset, OpenNeuro accession DS005700. No new participants were recruited. Informed consent and ethical approval were obtained by the original dataset creators; this re-analysis does not require additional ethical review.